

Introduction to Operating Systems

CPSC/ECE 3220 Spring 2020

Lecture Notes
OSPP Chapter 8 – Part B

(adapted by Mark Smotherman from publisher's
web site)

Efficient Address Translation

- Translation Lookaside Buffer (TLB)
 - Cache of recent virtual page -> physical page translations with permissions, in hardware
 - If cache hit, use that translation
 - If cache miss, walk multi-level page table
- Caches are expensive
- Static on-chip memory
 - Performance gain should be worth it
 - TLB lookup needs to be faster than full address translation
 - Some use multiulevel TLBs

What is Virtual Memory?

<https://www.youtube.com/watch?v=2quKyPnUShQ>

Cost of address Translation

Cost of translation =

Cost of TLB lookup (whether in TLB or not) +

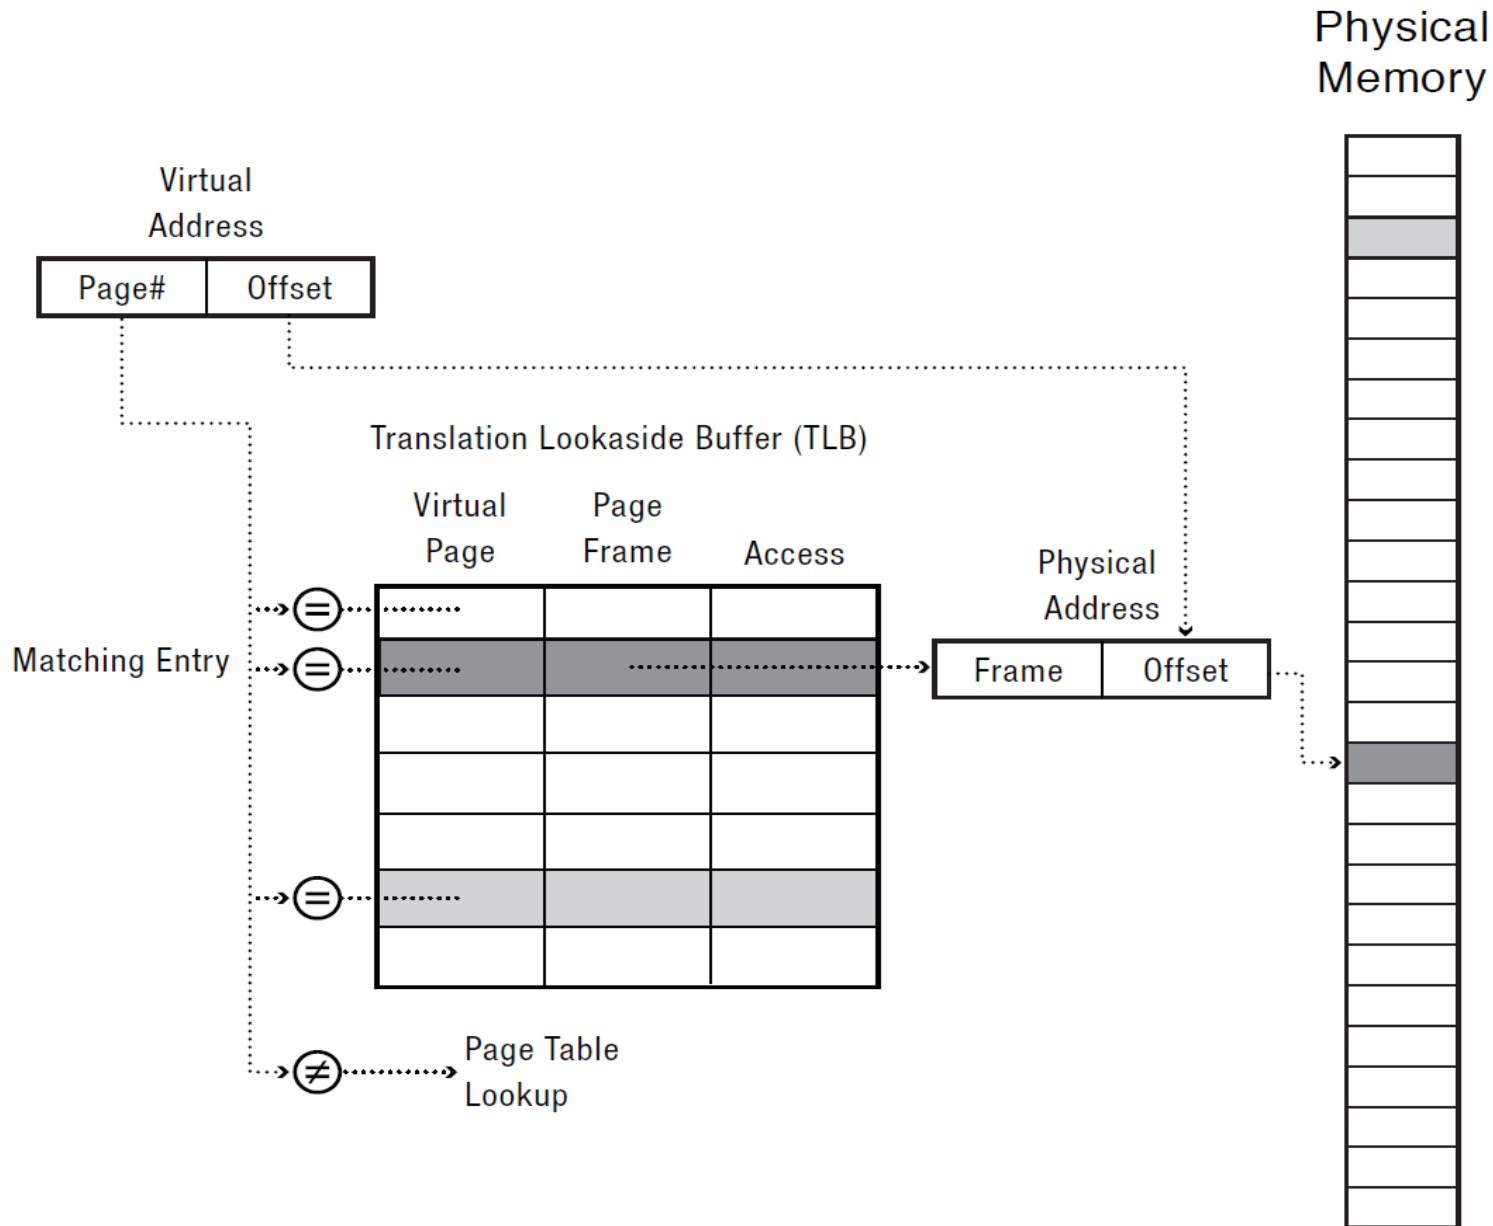
Cost of full translation (if not in TLB) *

Prob(TLB miss)

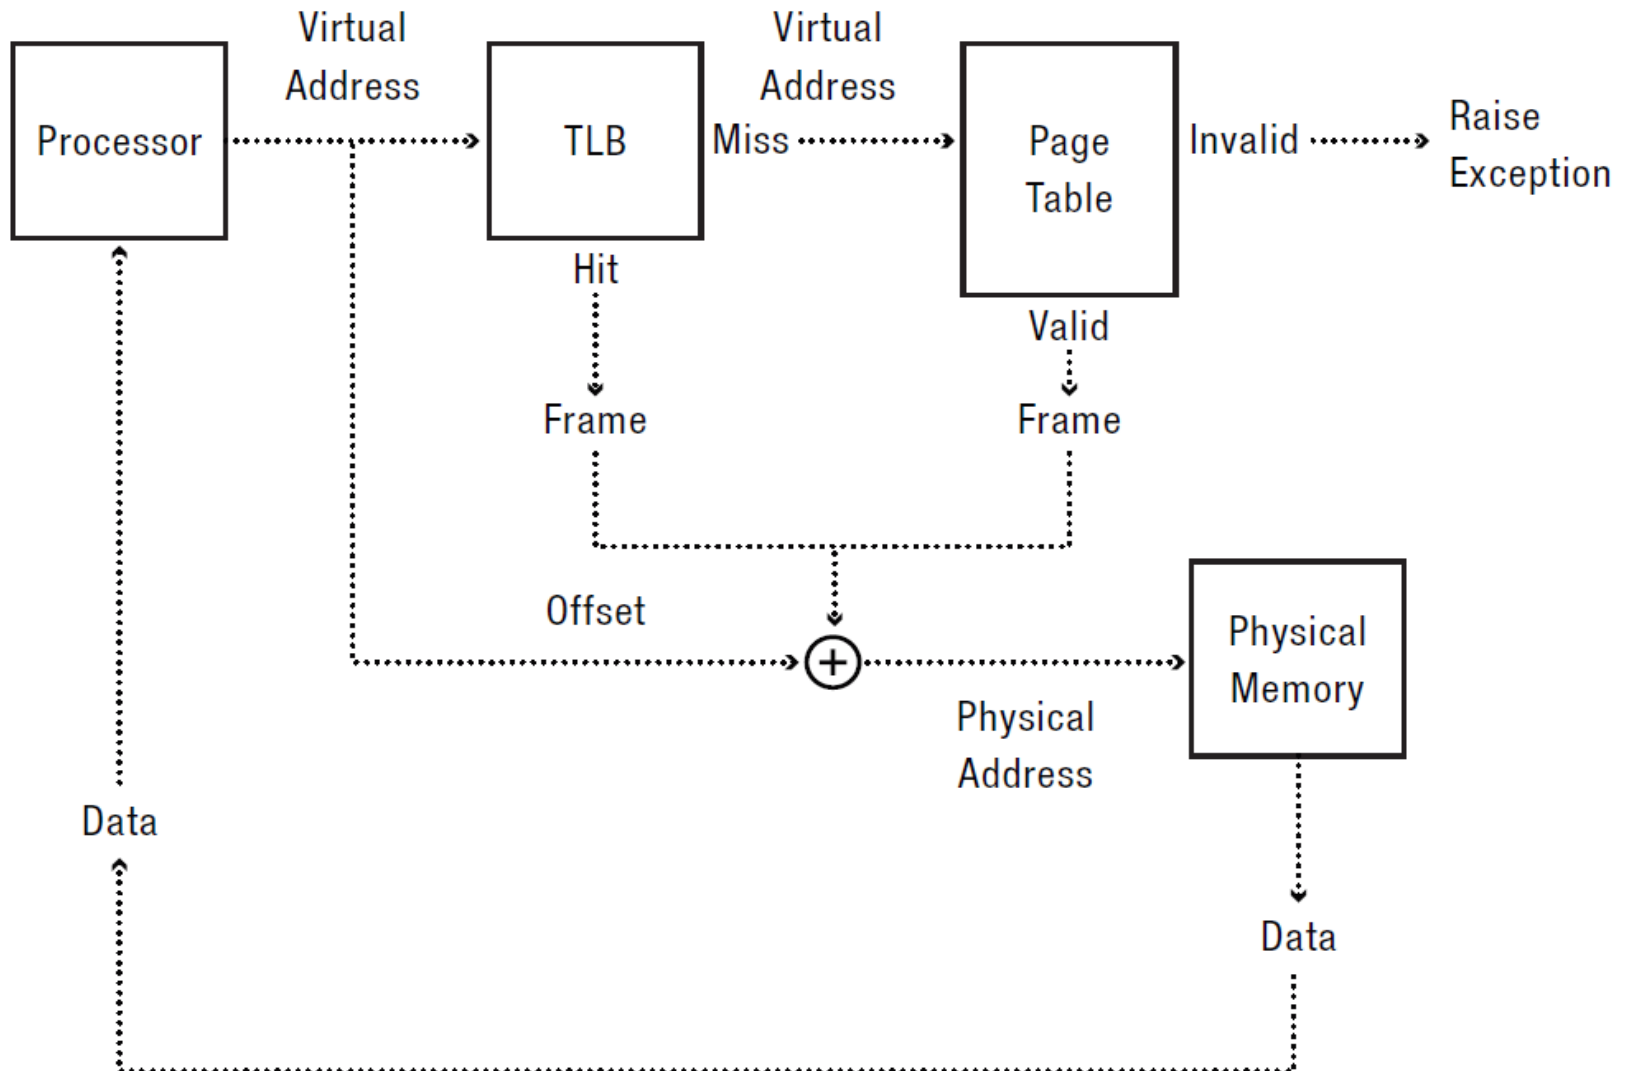
$\text{Prob(TLB miss)} = 1 - \text{Prob(TLB hit)}$

Design consideration: full address translation should be a rare event -> TLB should be sufficiently large

TLB Lookup



TLB and Page Table



Hardware Design Principle

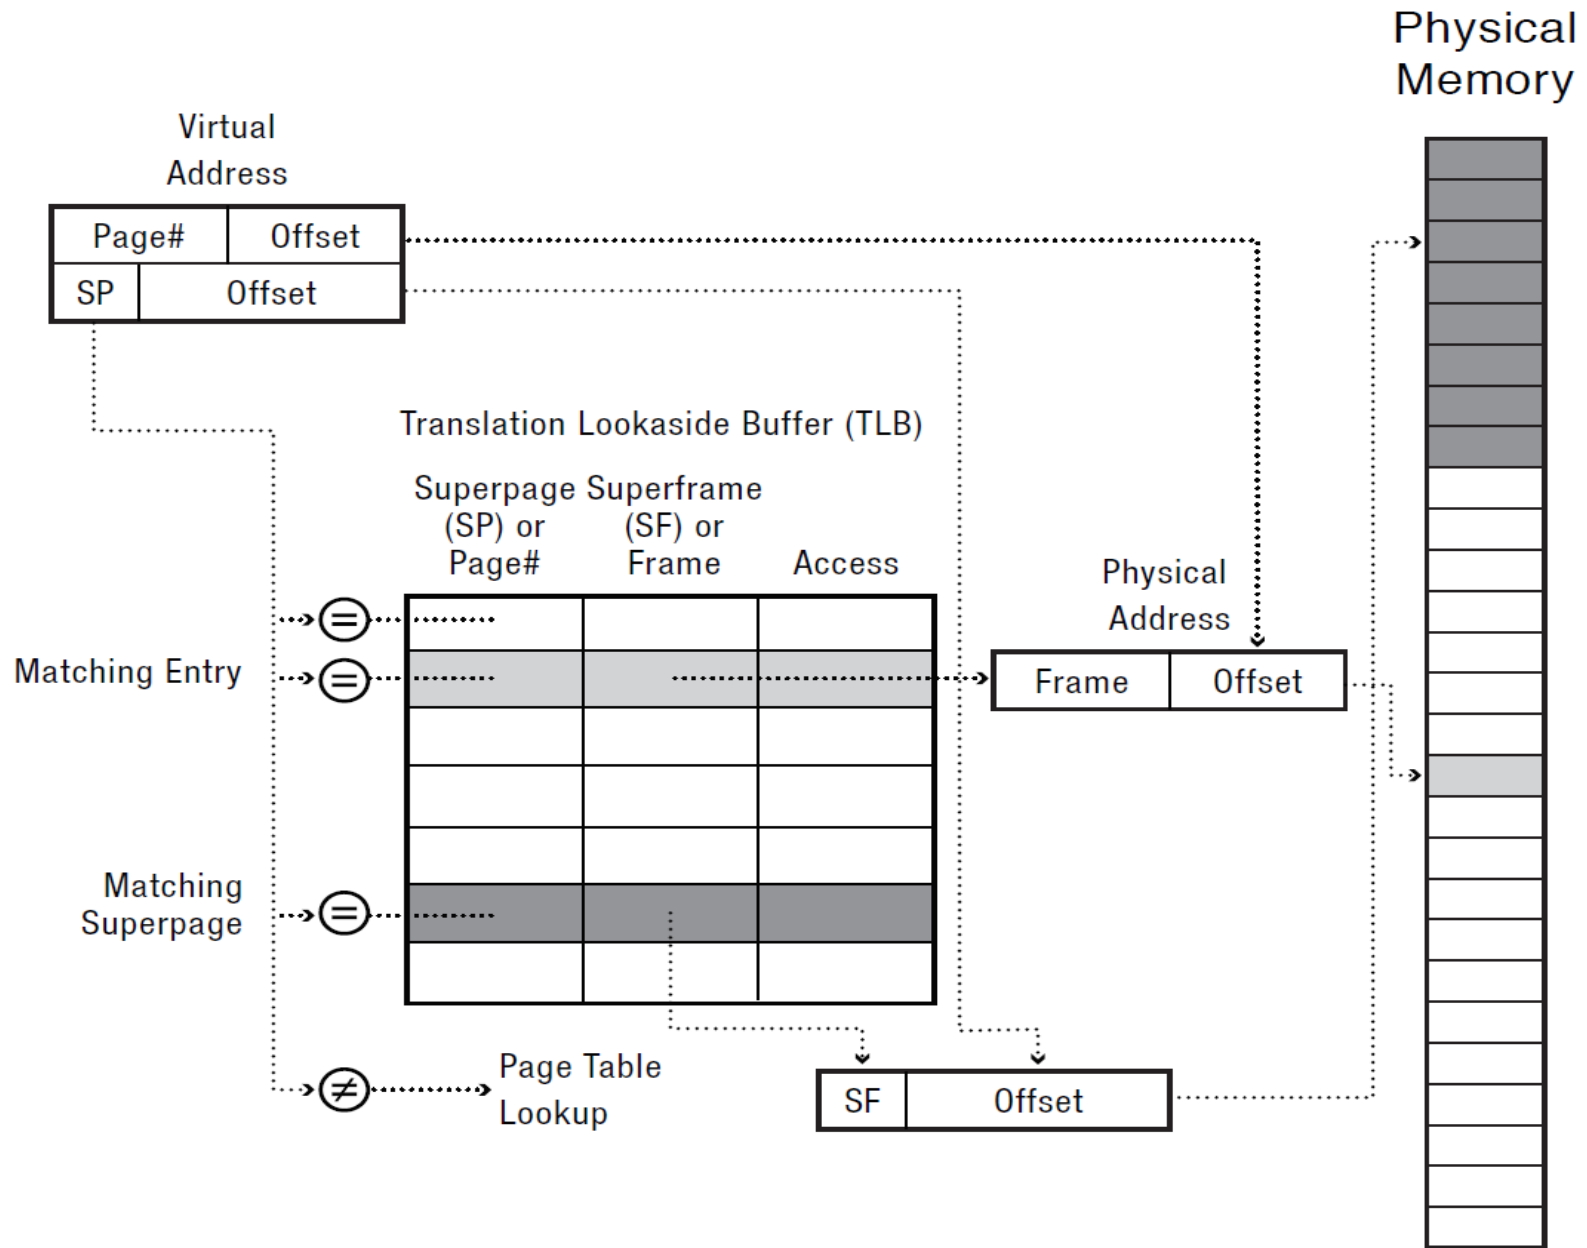
The bigger the memory, the slower the memory.

Why?

Superpages

- Improve TLB hit rate
- On many systems, TLB entry can be
 - A page
 - A superpage:
 - a set of contiguous pages in VM that map to contiguous pages in PM, aligned to share same high order address
 - TLB will have a flag to indicate a superpage
- Superpages are at the discretion of the OS.
- x86: superpage is a set of pages in one page table
 - x86 TLB entries
 - 4KB, 2MB, 1GB

Superpages



TLBs Consistency

- TLB needs to be in the consistent state
- TLB may contain old translation at context switch
- On every context switch – change hw **page table register** to point to new process' page table
- And *flush* TLB entry
 - Carries penalty

Tagged TLBs

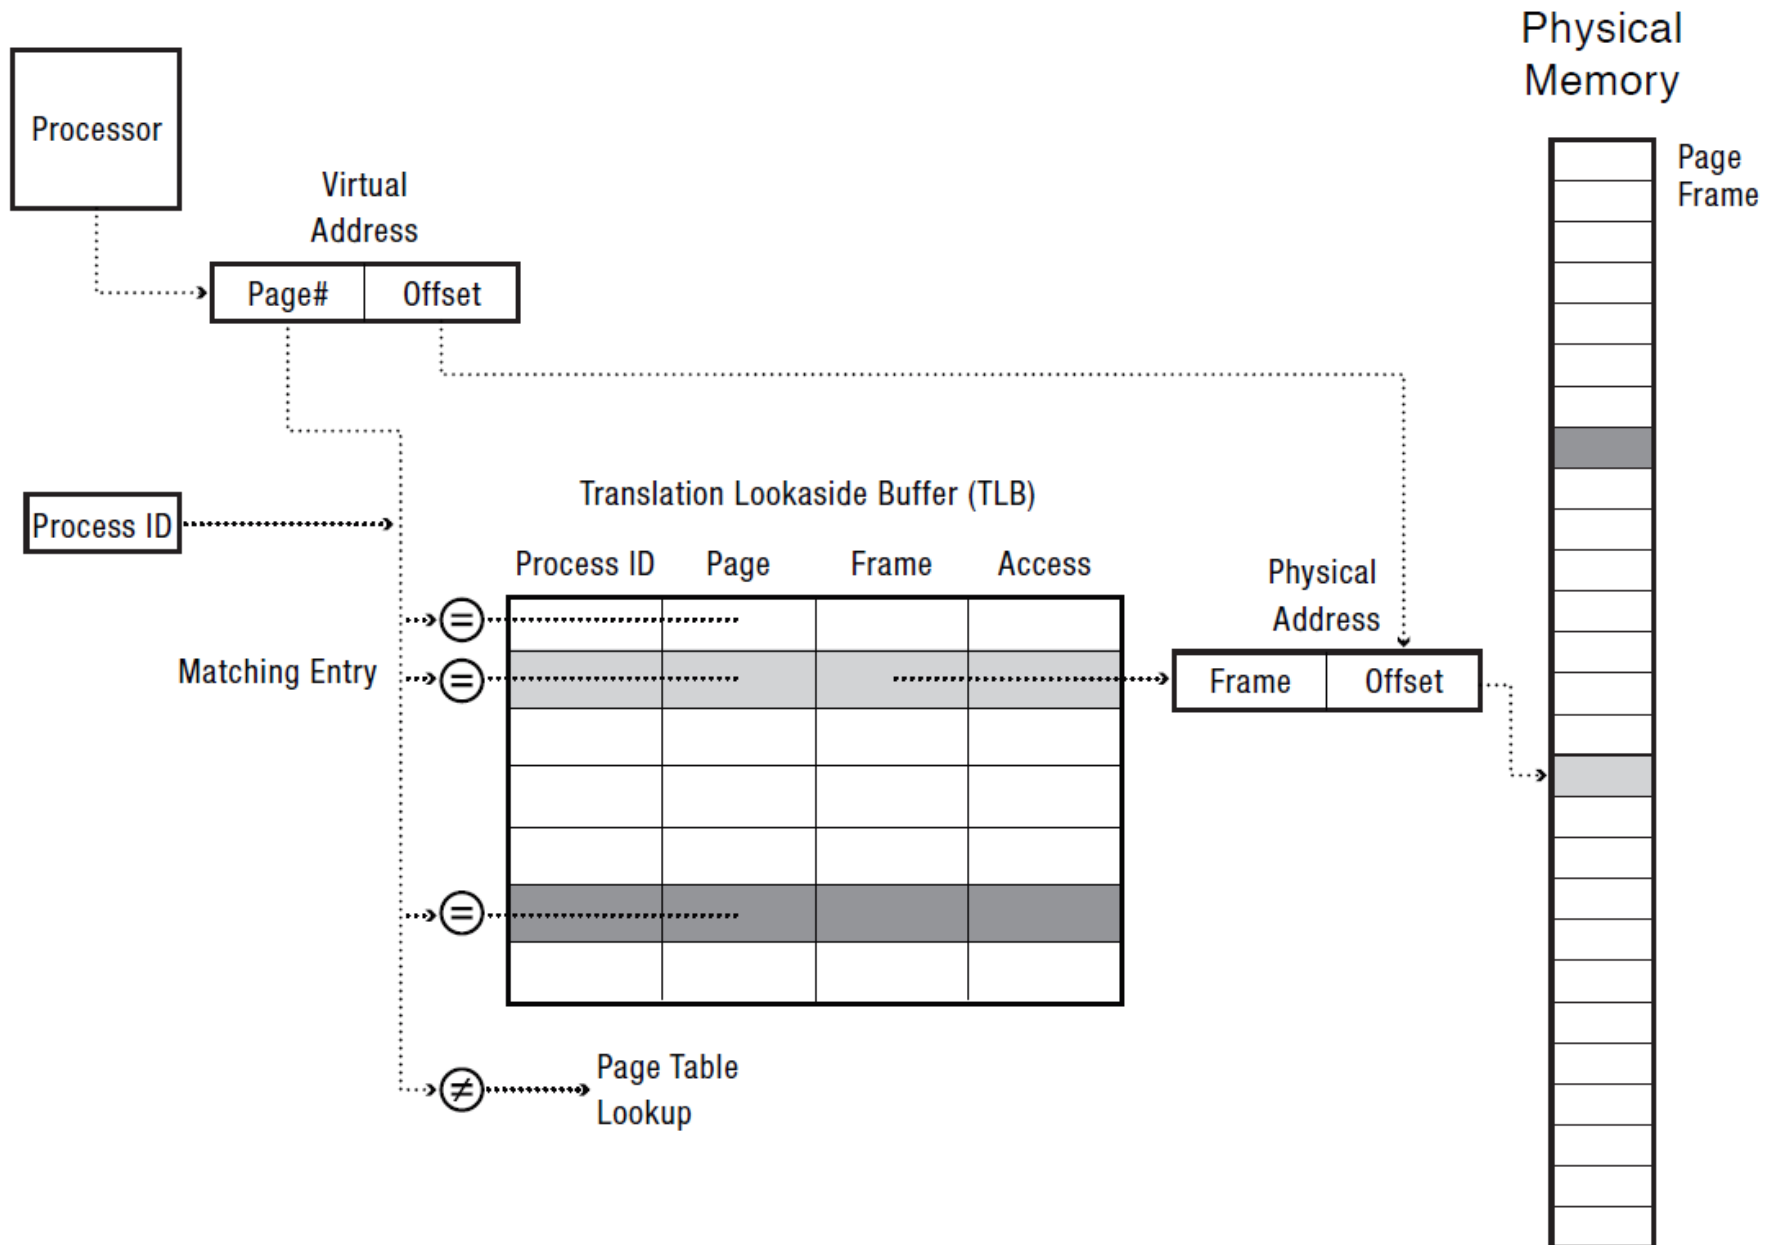
- What happens on a context switch?
 - Reuse TLB?
 - Discard TLB?
- Solution: Tagged TLB
 - Each TLB entry has process ID
 - TLB hit only if process ID matches current process

Multiprocessor TLB

- Any or all processors may have a cached copy of TLB
- Every time entry is modified, need to remove TLB from every processor
- OS interrupts every processor to remove old TLB entry - TLB shutdown
- Original processor continue executing *after* other processors remove entry

TLB Shutdown

		Process ID	VirtualPage	PageFrame	Access
Processor 1 TLB	=	0	0x0053	0x0003	R/W
	=	1	0x40FF	0x0012	R/W
Processor 2 TLB	=	0	0x0053	0x0003	R/W
	=	0	0x0001	0x0005	Read
Processor 3 TLB	=	1	0x40FF	0x0012	R/W
	=	0	0x0001	0x0005	Read



Address Translation Uses

- Process isolation
 - Keep a process from touching anyone else's memory, or the kernel's
- Efficient interprocess communication
 - Shared regions of memory between processes
- Shared code segments
 - E.g., common libraries used by many different programs
- Program initialization
 - Start running a program before it is entirely in memory
- Dynamic memory allocation
 - Allocate and initialize stack/heap pages on demand

Address Translation Uses (2)

- Program debugging
 - Data breakpoints when address is accessed
- Memory mapped files
 - Access file data using load/store instructions
- Demand-paged virtual memory
 - Illusion of near-infinite memory, backed by disk or memory on other machines
- Zero-copy I/O
 - Directly from I/O device into/out of user memory
- Efficient support of virtual machines

Address Translation Uses (3)

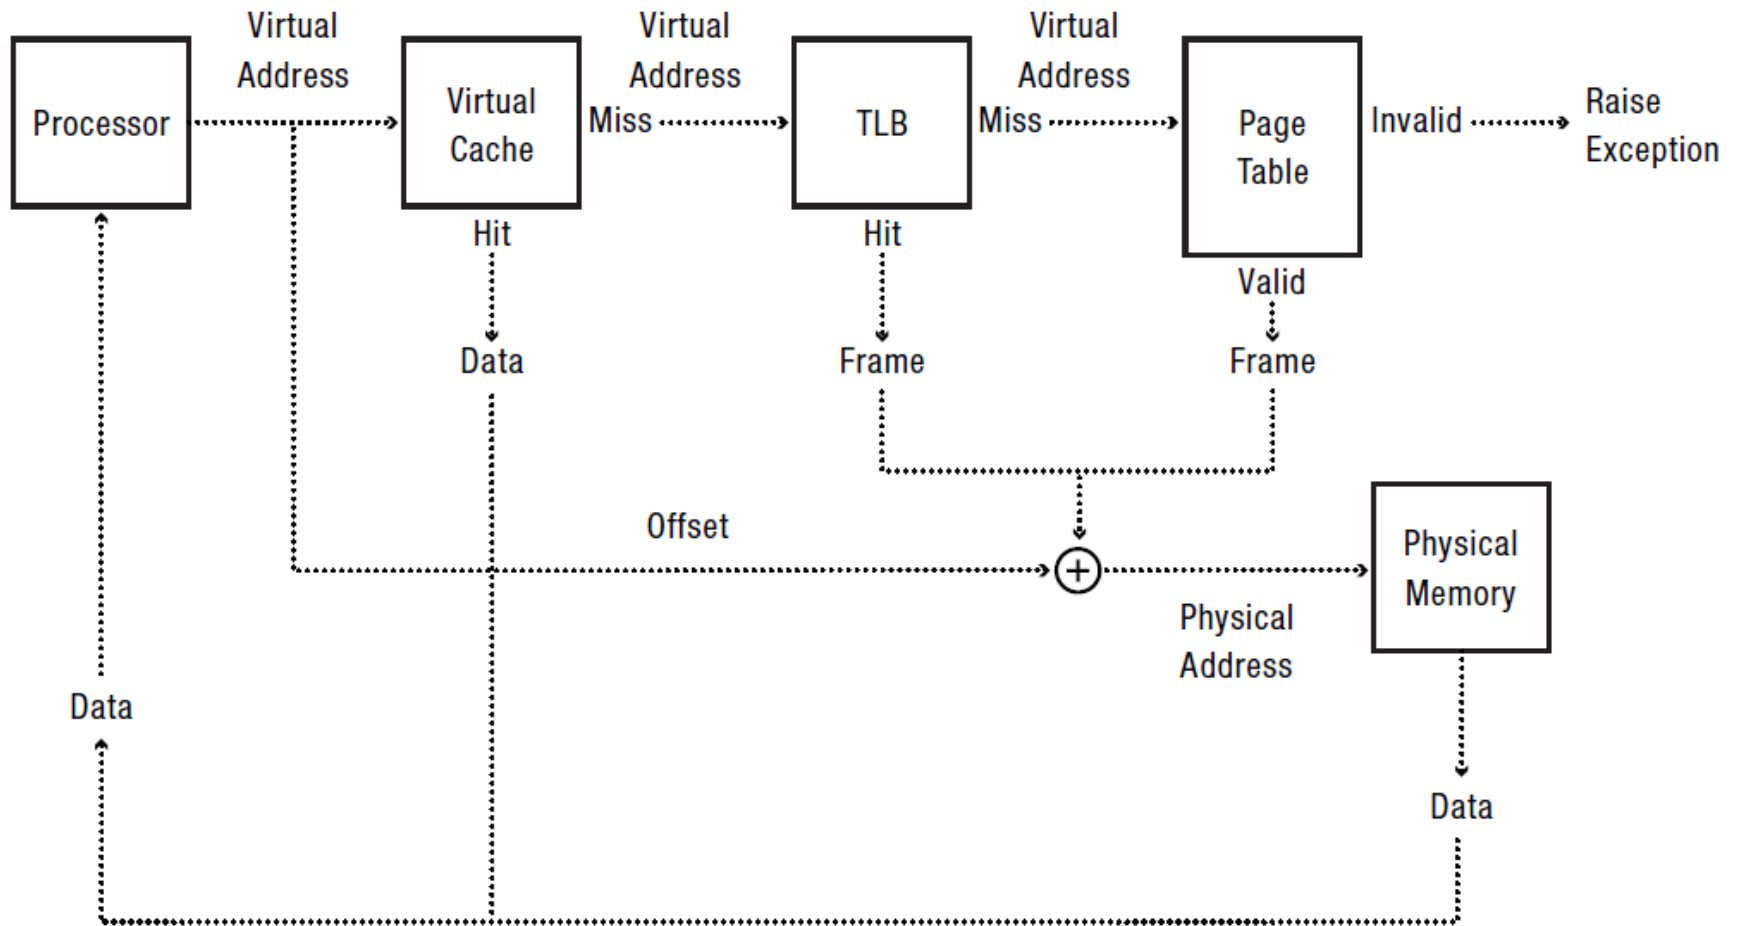
- Checkpointing/restart
 - Transparently save a copy of a process, without stopping the program while the save happens
- Persistent data structures
 - Implement data structures that can survive system reboots
- Process migration
 - Transparently move processes between machines
- Information flow control
 - Track what data is being shared externally
- Distributed shared memory
 - Illusion of memory that is shared between machines

(if time permits)

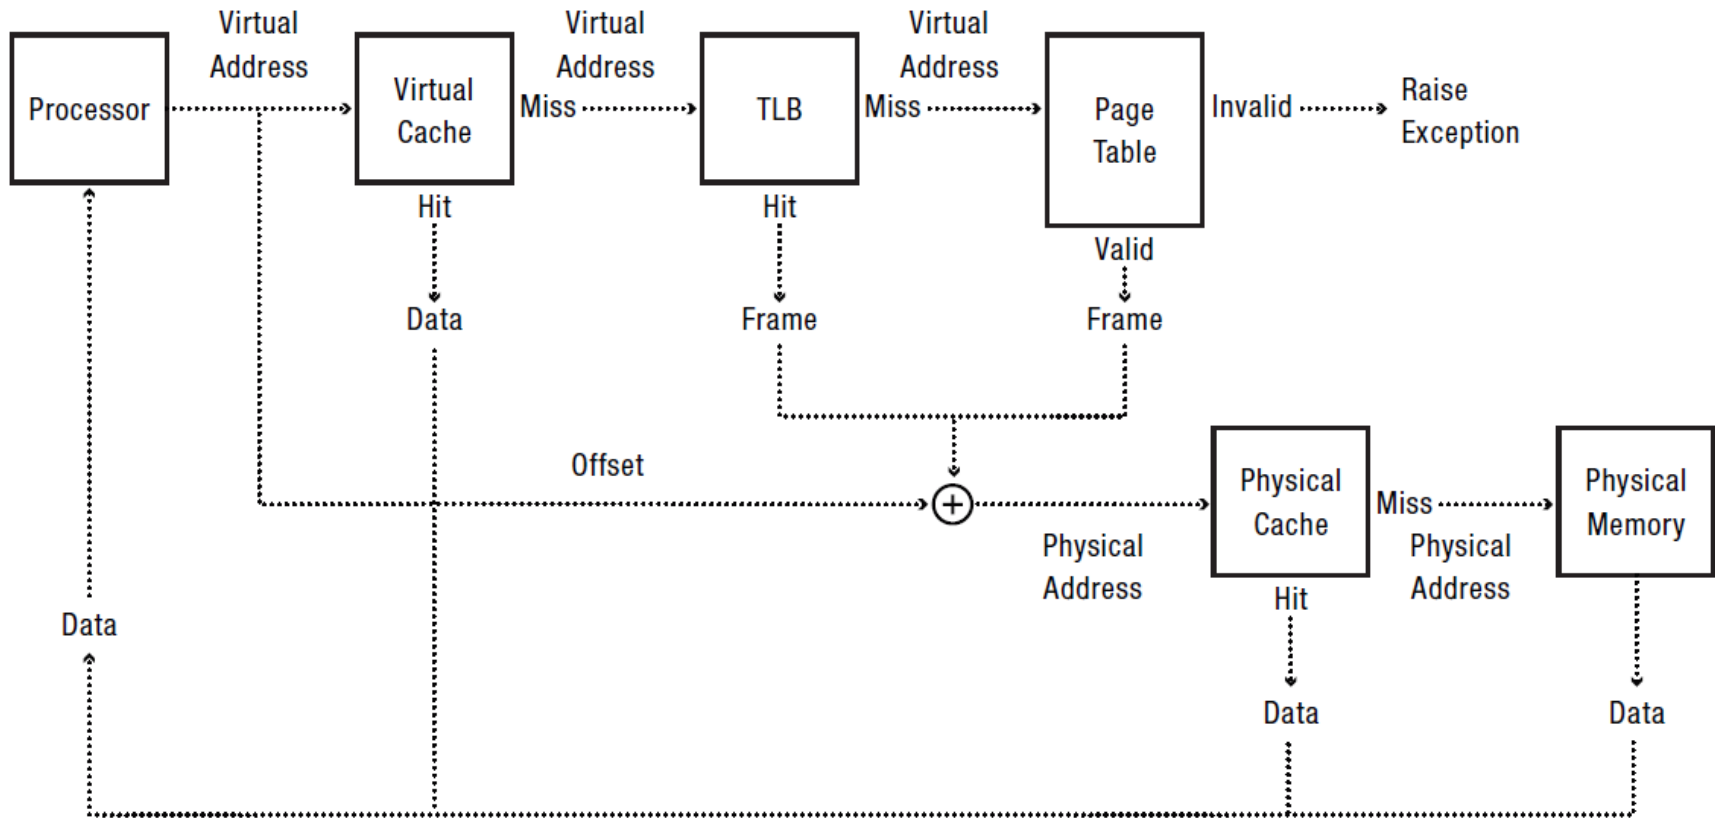
Virtually Addressed vs. Physically Addressed Caches

- Too slow to first access TLB to find physical address, then look up address in the cache
- Instead, first level cache is virtually addressed
- In parallel, access TLB to generate physical address in case of a cache miss

Virtually Addressed Caches



Physically Addressed Cache



Question

- With a virtual cache, what do we need to do on a context switch?