

Threads

September 9th

Why processes

- Sharing
- isolation
- concurrency

Threads act like processes (light weight processes)

Concurrency vs Parallel
Switches tasks rapidly

Runs processes at some time

Threads

Kernel threads - kernel manages threads

User threads - created and managed by user

User mode benefits

- Super fast, full scheduler controller
- Cons - Safety concerns, no parallelism

You cannot predict how long for a process to run

Why? Multiple other things running + the scheduler may not prioritize your requesting

Creating Threads in C

pthread_t t1, t2;

[
pthread_create(&t1, NULL, threadfunc, NULL);
↓

pthread_join(t1, NULL);

pthread_join(t2, NULL);

- Concurrent race condition - Depending which process runs which code changes output

Mutex - locks

pthread_mutex_t * thelock = PTHREAD_MUTEX_INITIALIZER

front() { pthread_mutex_lock(*thelock);

i++;

pthread_mutex_unlock(*thelock); }

← Called Critical Section

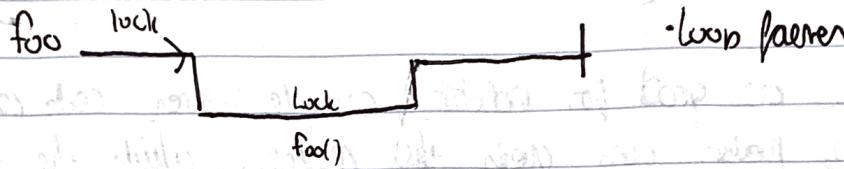
Extremely slow involves a loop

In order to cancel a thread (pthread_cancel)
 the thread has to be in a canceled state (sleep is a canceled thread)

* Attributes (new in updatons)

4hr 40 min

Thread Safety - if you can run code in multithread and it wakes threads
 Reentrant function can `foo();` interrupt `foo()` at any point
 then returns `foo()` return, then resume the original `foo()`, its previous execution



Printf calls on malloc memory contain some information
 to move that the memory come from printf
 \rightarrow `printf("Name");`

maybe memory has "none" inside of it its allocated memory
 since we stored pointers to malloc

OOSP CH 3.1 - 3.2

Create Process

- Initialize process control block
- Initialize new address space
- Load program into address space
- Copy arguments into memory in address space
- Initialize hardware context to start creation of thread
- Inform scheduler process is ready

Unmap Exec

- Load program into current address space
- Copy arguments into memory address space
- Initialize hardware context to begin of thread

our Scheduling Algorithm

PID	ARR	Length	Deadline
1	0	3	8
2	2	2	7
3	3	5	9
4	5	1	6

- Ratio of L/DL
- EDF (earliest deadline first)

Using EDF (4213)

- [Can run into starvation if deadlines are too tight]

~~Execution Context~~
~~Least Slack~~

- How many quanta can I waste while completing the task?

Chapter 4 Concurrency and Threads

4.1 Thread use cases

- Helpful for games when you sharing graphics, receiving input, and updating player location simultaneously

4.2 Thread Application

thread - a single execution sequence that represents a separately schedulable task

- Can run suspend or resume at any time

Threads are virtual machines with immeasurable speed

thread join is analogous to wait for fork and exec

thread - state does generally the same things fork does + exec

4.4 Threads data structures and life cycle

threads have a per thread state and shared thread state

- Thread control block stores information about threads executing and meta data used to manage them
- Each thread also has its own stack which is pointed to by TCB
- Thread meta data contains information such as thread ID, scheduling priority, and status

Thread data should be a linked list

- Shaded state - state shared between running threads
- Program code is shared between threads to a extent.
- Information visible to each thread is the ~~process~~ heap, and global variables

4.5 Thread Life Cycle

Initialize \rightarrow Ready \rightarrow Wait \rightarrow Run \rightarrow Finish

thread_create(), thread_join(), thread_yield(), thread_exit()

Initialization places thread into init state, allocates per thread information, puts thread into ready state by placing threads onto ready priority queue

Ready - At any time, the threads repository live on thread control block, can be taken off and copied to CPU repository while running

Running - Thread is executing on CPU, registers are copied to CPU, can go back to ready by yielding, or copying registers back to thread control block, and switching CPU to run next thread

* A thread running is not on the ready list and vice versa
 Wait - thread is waiting for some event. Cannot run until some thread finishes and changes its state from waiting to running
 Thread is stored on waiting list of some synchronization variable associated

Finished - Finished thread can never run again. OS can reclaim memory, but sometimes thread is placed on a ~~finished~~ ^{finished} list. Once the exit value is read by thread_join, the thread is no longer used

4.6 Kernel Threads

3 Step Creation

1. Allocate per thread state 3. Put TCB onto ready list

2. Initialize per thread state

- thread starts in dummy function, which then calls derived function, more function returns, when which stub returns to

whatever function at top of stack

Deleting Thread

- Remove thread from ready list
- Free per-thread state allocated for thread

* Interrupt could stop thread mid calculation and cause memory leak

* Thread cannot see its own state!

Instead, copy thread to finished state, one another thread has a moment clean the memory of threads in finished state

Thread Context Switch - Suspends current thread and resumes some other thread

- Saves current ^{running} thread registers to TCB and stack, restores new thread information to CPU registers

(caused by):

- Call to actual thread library
- CPU interrupt or memory exception

Threads should not know if interrupt or memory exception occurred.

To the thread this won't happen. That's why thread running order is very unpredictable

Must disable interrupts during switch or bugs will occur

* Page 62 contains useful explanation

Involuntary kernel Thread Context Switch

- Interrupt or trap signifies kernel thread switch
- Save state
- Run kernel thread
- Restore state

* Review important

Thread Safe Redefines

- Global shared memory
- Static values

4.7 Continued 4.8

Threads in a process contain

- User level stack
- Kernel interrupt stack for when thread comes interrupt
- Kernel TCB for saving and restoring per thread state

Scheduler activations - user level threads scheduler is notified for every kernel event that might affect user level threads

Asynchronous IO - much more practice in server applications

Chapter 5. Synchronizing Access to Shared Objects

Independent threads - threads operate on completely separate subsets of memory

Cooperating threads - read and write shared memory

Free condition - thread output depends on interleaving operations of different threads

Atomic operations - indivisible operations that cannot be interleaved or split by operations

The challenge with threads accessing shared blocks of memory is it's extremely easy to modify something that should not have been changed. Such as the too much milk problem. Someone checks the fridge, no milk, heads to the store and by person buys the exact same thing, and buys milk. The solution is to throw a lock around the particular task it's handling. If they when you complete the task you pass it off to the next thread or AKA unlock it.

Shared objects - objects that can be safely accessed by threads

5.2

Synchronization Variable - Data structure used for coordinating access to a shared state

- AKA locks and conditional Variables

Locks 5.3

- Synchronization Variable that provides mutual exclusion. When one thread holds a lock no other thread can.

Locks guard a shared Objects state

- Generally a lock for each public method

Locks can be in one of two states

- Busy or free
- Lock is always initially free
- ~~if Acquire()~~ waits for lock to not be busy and ~~if Release()~~ waits for lock to be released

Lock should ensure following properties

- At most one thread holds the lock
- if no thread has it lock, a thread ceiling acquire should eventually acquire the lock
- bounded waiting - If asks for the lock, then there is a bound on how many times a thread can polls the lock before T

5.3.2 Thread Safe Queues

(Critical Section) - area of shared memory that must be accessed atomically

- Threads should only share dynamically allocated memory

5.4 Condition Variables

- Provide a way for one thread to wait for another thread to take action
 - Very helpful when thread is waiting for a lock to open up.

wait(lock * lock) - releases lock of calling thread, and places thread on waitlist. When execution resumes is enabling, lock is re-acquired before returning from wait

signal() - removes one thread and marks as ready to run to the scheduler. No effect if no threads are present

broadcast() - Takes all ~~threads~~ CV's waiting lists and marks them as ready to run

* Page 130 is useful

- No threads can wait unless the thread has the lock in ~~possession~~
- Wait must always be called inside a loop

Thread Life Cycle

- Running - thread that calls wait is placed on waiting list
 - copy TCB from running to waiting block
 - calling thread copies TCB to ready state, and the OS schedules the makes it run

Generally you have as many conditional variables as unique ways ~~and~~ in which a method must ~~work~~

Left off on page 139.

Project 2 Notes

User Mode

There is cooperative and preemptive mode

Cooperative Mode

Emulate interrupts or autoquakes

Context Switches

- Save the state of a thread before switching
`(get context)(First-context)`
 - Save current context
`set context(First-context)`
 - Load current context
`swap context(second, First)`
 - Save one, load the other
- * `myContext.uc-stack.ss-sp = a_new_stack`
- * `myContext.uc-stack.ss-size = STACK-SIZE`

`make context(&myContext, void (*)() thread_func, 1, 7);`

So, when I call

`swap-context(&main-context, &thread);`

- Saves current context into main-context and loads the current context of thread which ~~func~~ points to the function `thread-func`

The addresse should be a linked list node or array
where function ptr into another function, that can swap
the contexts when needed?

This is a linked list storing each threads context

Threads

Use Case :

Threads are useful when you want multiple things occurring at once. You want to render a scene while allowing the user to use a responsive main menu.

Or rendering generation while still allowing the player to move around.

Thread Abstraction

- Basically threads act like mini virtual machines running simultaneously tasks. Allows that how they appear. Instead threads execute a sequence of instructions until a switch occurs. Then that thread is essentially "popped" while that thread is popped another one takes over and keeps working on another section until its switched off. Threads give the illusion of everything simultaneously being worked on.

join
yield

Thread API

join : Initializes a thread and joins it together with the desired function to run?

wait : Takes current running thread, removes context, lets next thread begin running. Once wait returns the popped thread resumes running

yield : Thread gives up running on the machine to let another thread run. At any time the thread may be signaled to run again. But until then thread is on waiting list

Thread yield

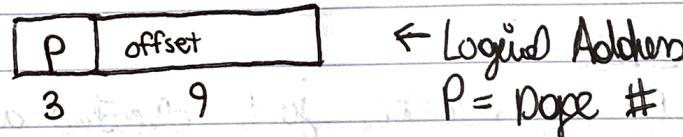
- Need idea of currently running thread
- (will) have thread pointer, but at init that is known the threads that point to running thread

Memory

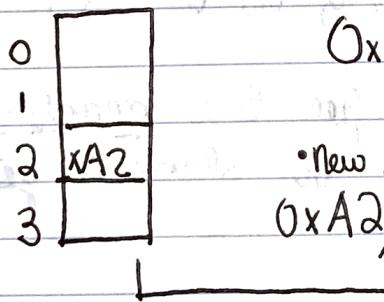
External Fragmentation - The memory heap covered has a bunch of small memory pieces that are too small to satisfy the memory request

Ways to avoid fragmentation

- FIFO Fit - Find the block that first fits
 - Next Fit -
 - Best Fit - Find smallest space it can fit into
 - Worst Fit - Largest place it can fit into
- Instead of variable sized blocks, the O.S. uses fixed sized blocks
- Pages - Some sized chunks in logical address
Frames - Some sized chunks in physical address space



Example



• New Address