

# Introduction to Operating Systems

CPSC/ECE 3220 Spring 2020

Lecture Notes  
OSPP Chapter 11

(adapted from publisher's slides)

# File Systems

- Abstraction on top of persistent storage
  - Magnetic disk
  - Flash memory (e.g., USB flash drive)
- Devices provide
  - Storage that (usually) survives across machine crashes
  - Block level (random) access
  - Large capacity at low cost
  - Relatively slow performance
    - Magnetic disk read takes 10-20M processor instructions

# File System as Illusionist: Hide Limitations of Physical Storage

- Persistence of data stored in file system even if :
  - crash happens during an update, disk block becomes corrupted, flash memory wears out
- Controlled access to shared data
- Naming:
  - Named data instead of disk block numbers
  - Directories instead of flat storage
- Performance:
  - Data placement and data structure organization
  - Cached data

# File System Abstraction

- File system
  - Persistent, named data
  - Operating system crashes (and disk errors) leave file system in a valid state
  - Journaling feature
- Access control on data
- File: named collection of data
  - Hierarchical organization (directories, subdirectories)
- Performance
  - Achieve close to the hardware limit in the average case

# Files

- Metadata
  - owner, access permissions, timestamps (creation, last written), size, reference count, lock, etc.
- Data
  - May be unstructured or structured:
    - Stream of bytes (even if stored as blocks)
    - Records are collections of related fields, often with a key field used for searching and sorting

# Identifying the File Type

- Metadata
- Extension identifier (.c, .o)
  - Can be used to specify usage or structure
  - OS can associate a specific application with an extension
  - Not all OSs require extensions
    - Does Linux require extensions?
- Magic number within file
  - Identify usage or structure

# Three File Types in UNIX

- Regular file
  - Text (e.g., ASCII)
  - Binary (e.g., executable)
- Directory
  - Writes must be restricted to preserve structure
- Special file
  - Maps physical I/O device to the file system
  - Linux device files

# File Type Philosophies

- UNIX philosophy as stated by Ritchie and Thompson (inventors):
  - “the structure of files is controlled by **the programs that use them**, not by the system”
  - Why do you think this is so?

.



# Directories

- Directory
  - Group of named files or subdirectories
  - Mapping from file name to file metadata location
- Path
  - String that uniquely identifies file or directory
    - Ex: /web/home/joe/public\_html/3220.html
- Links
  - Hard link: link from name to metadata location
  - Soft link: link from name to alternate name (which links to metadata location)
- Mount points
  - Mapping from name in one file system to root of another

# Path Names

- Absolute path
  - Fully qualified
  - Starts at root (“/” in Unix systems, “C:” in Windows)
- Relative path
  - Partially qualified, ex: public\_html/3220.html
  - Unqualified, ex: 3220.html
  - Starts at current working directory
  - . => this directory, .. => parent directory

# UNIX File System API

- create, link, unlink, createdir, rmdir
  - Create file, link to file, remove link
  - Create directory, remove directory
- open, close, read, write, seek
  - Open/close a file for reading/writing
  - Seek resets current position
- fsync
  - File modifications can be cached
  - fsync forces modifications to disk (like a memory barrier)

# File System Workload

- How are files used?
  - Most files are read/written sequentially
  - Some files are read/written randomly
    - Ex: database files, swap files
  - Some files have a pre-defined size at creation
  - Some files start small and grow over time
    - Ex: program stdout, system logs

# Connection-Oriented Interface

- Explicit open and close operations for files
  - OS creates an internal data structure on open
- Read and write ops use descriptor (a.k.a. handle or stream) to identify the internal data structure
  - No need to reparse file name
- Per-open data structure contains:
  - Access permission under which file was opened
  - Location of file (e.g., inode number)
  - Pointer to current byte or record for sequential access

# UNIX Data Structures for Open Files

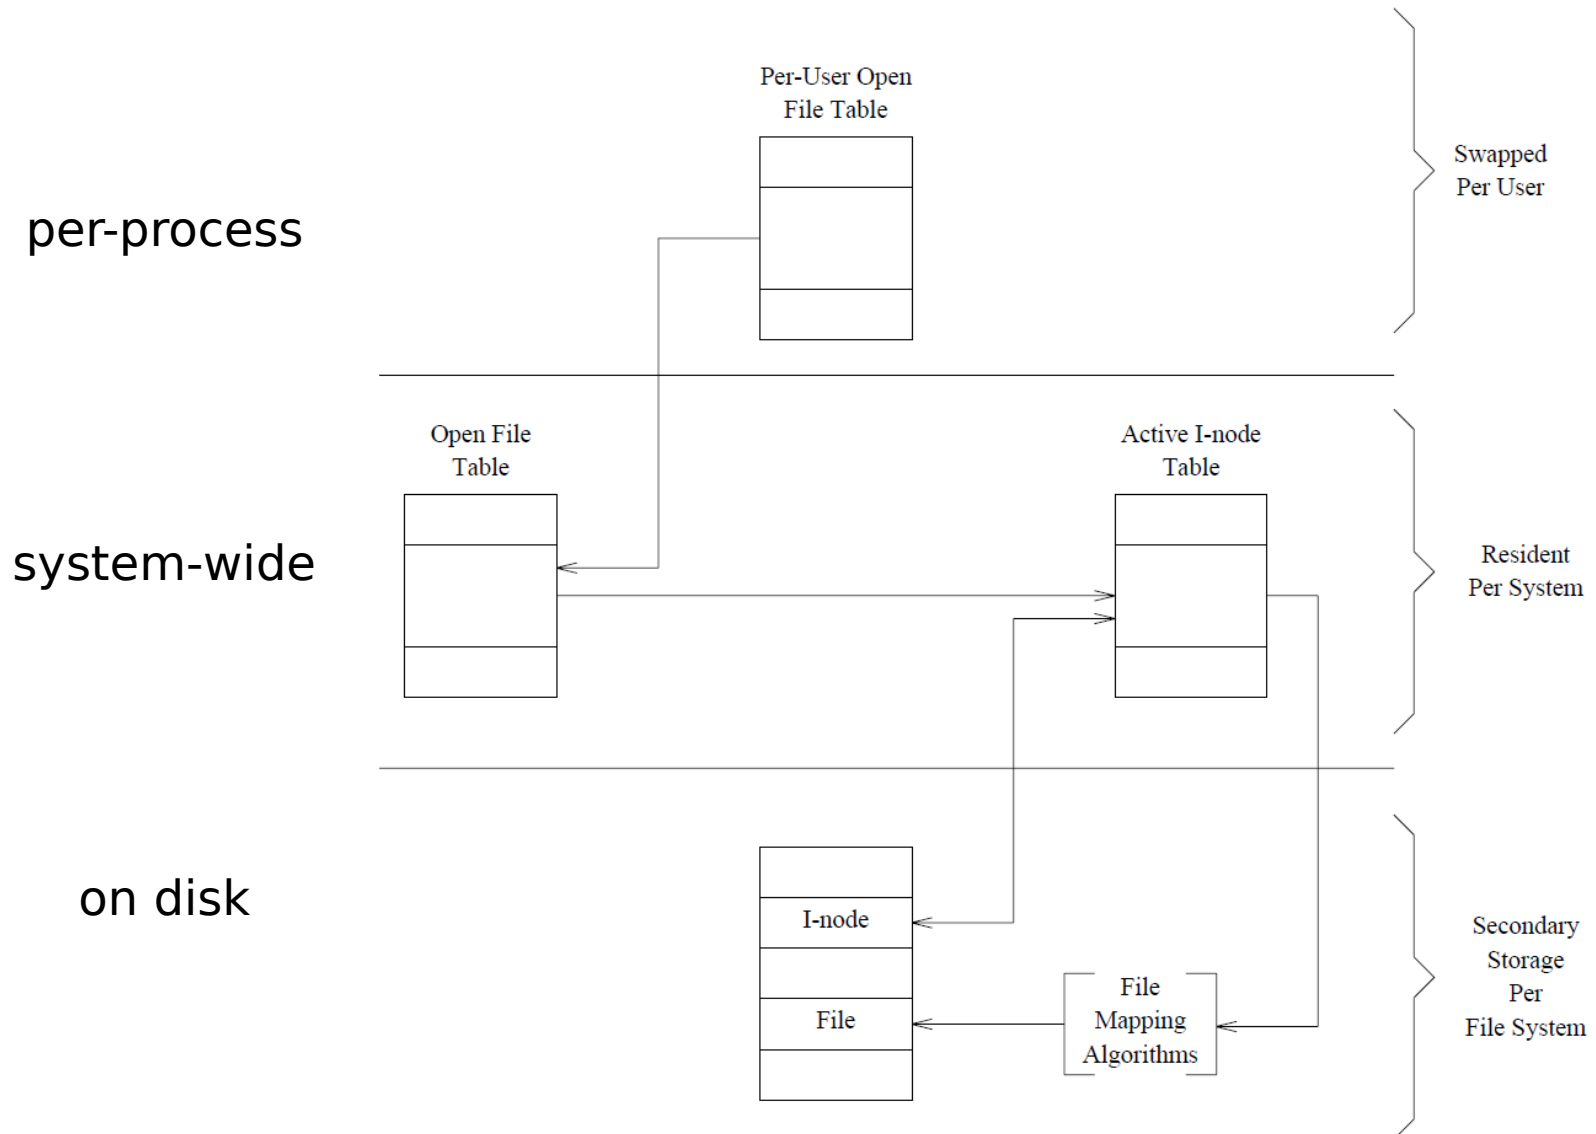
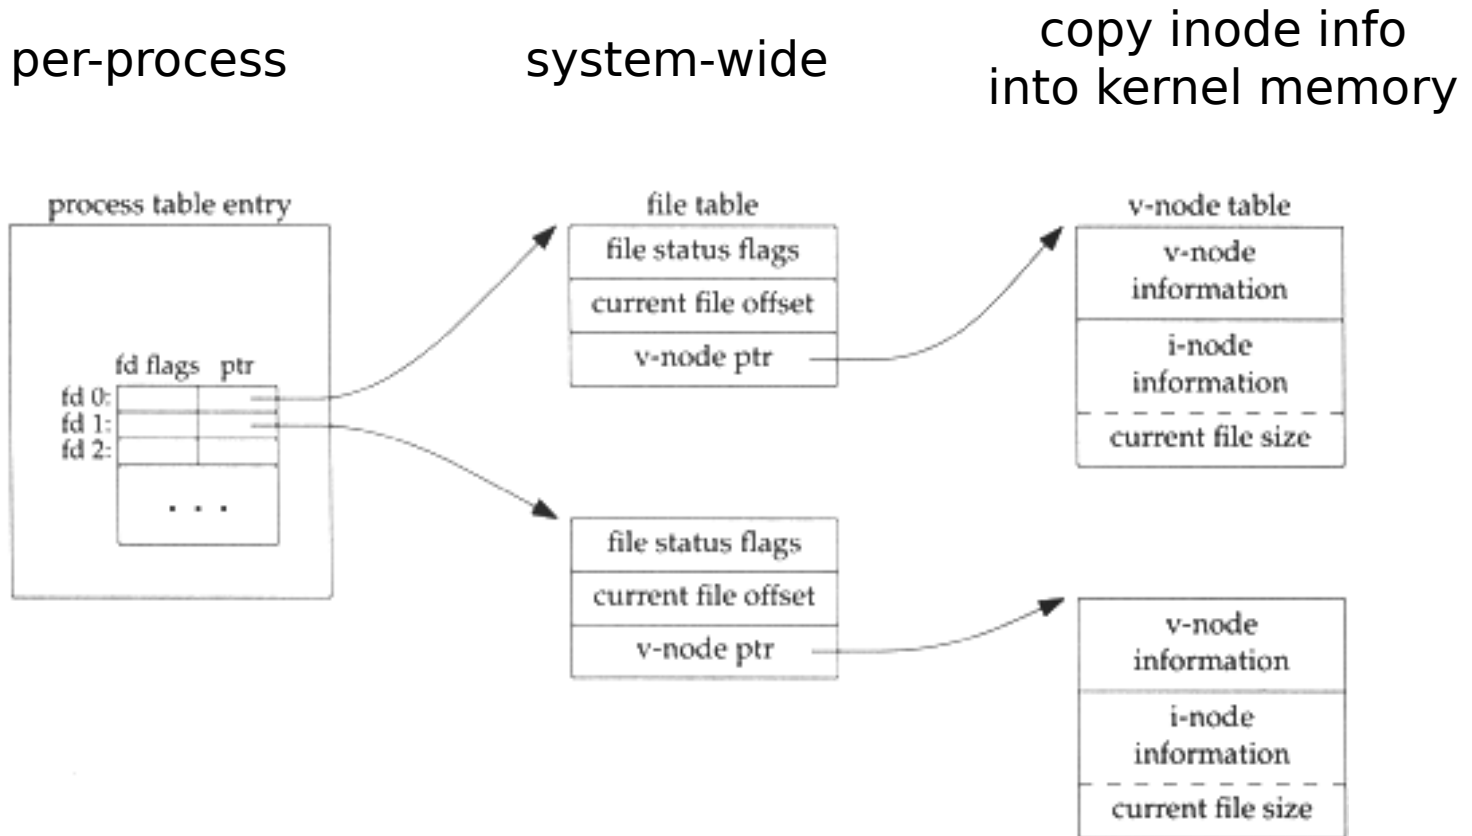


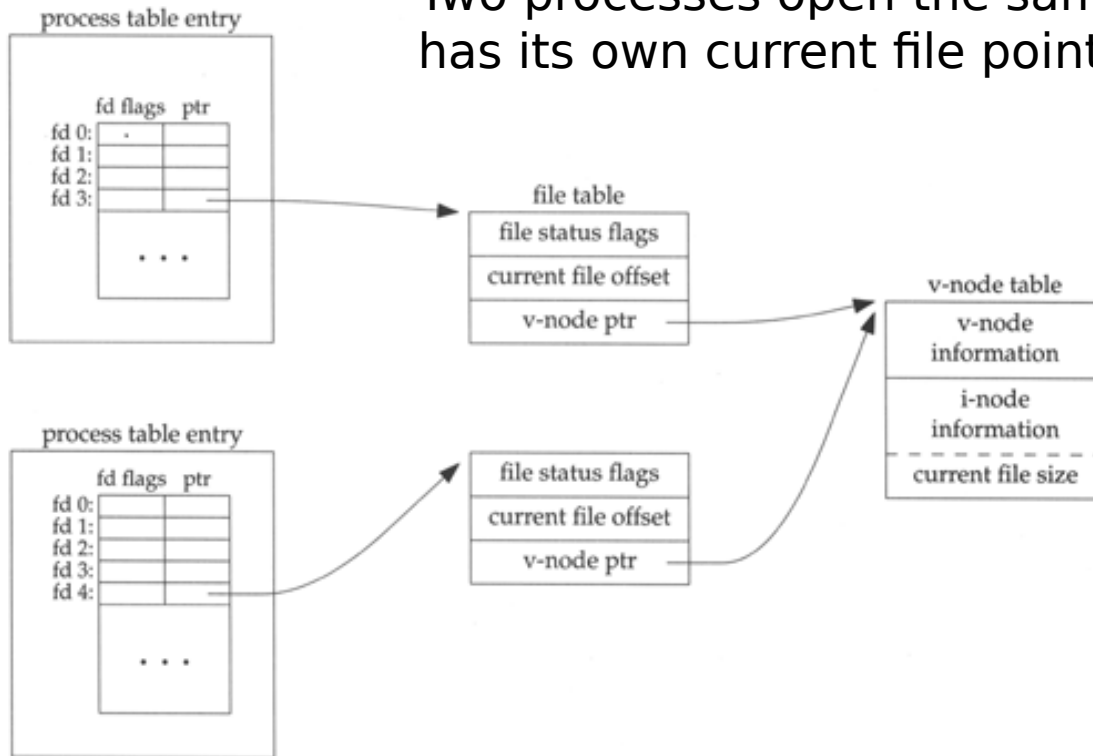
Diagram from K. Thompson, "UNIX Implementation," Bell System Tech. Jnl., 1978

# UNIX Data Structures for Open Files (2)



# UNIX Data Structures for Open Files (3)

Two processes open the same file – each has its own current file pointer (“offset”)





# File System Workload

- File sizes
  - Are most files small or large?
    - SMALL
  - Which accounts for more total storage: small or large files?
    - LARGE

# File System Workload

- File access
  - Are most accesses to small or large files?
    - SMALL
  - Which accounts for more total I/O bytes: small or large files?
    - LARGE