

Slides to Accompany *Programming Languages
and Methodologies*

R. J. Schalkoff

Chapter 15, Part 1: Parallel Programming

Parallel Programming: Algorithms and Implementation

- Key issues of parallel programming:
 - parallel algorithm decomposition
 - parallel hardware development (or availability)
- A very real and significant problem is *how to (perhaps automatically) decompose a given processing task into one which may be executed in parallel segments.*

Questions

1. Where is the parallelism in the desired computation?
2. How do I articulate/exploit it in my programming language?

The answers usually require some consideration of the underlying computational architecture.

A Simple minic Example

Consider the simple minic code fragment:

```
a=10;  
b=sqrt(a);  
c=a+b;  
d=6;  
e=sqrt(d);  
f=d+e;  
g=f+c;
```

Figure 1: Simple minic Code Fragment for Potential Parallelism Analysis

In this example, it is reasonably easy to (informally) identify parts of this code which are suitable for parallel execution. For example, two blocks:

```
a=10;  
b=sqrt(a);  
c=a+b;
```

```
d=6;  
e=sqrt(d);  
f=d+e;
```

could be computed concurrently. However, the computation required by

```
g=f+c;
```

must necessarily follow these computations.

Example: Control Constructs and Parallelism

Consider the code fragment:

```
b=sqrt(a);  
if (b<4) then  
    if (a==1) b=4;  
c=a+b;  
d=6;  
if (c>a) then d=1;  
e=sqrt(d);  
if (e>a) then e=1;  
f=d+e;  
g=f+c;
```

Figure 2: minic Code Fragment With Conditional Constructs

Clearly identifying parallelism in this case is more challenging.

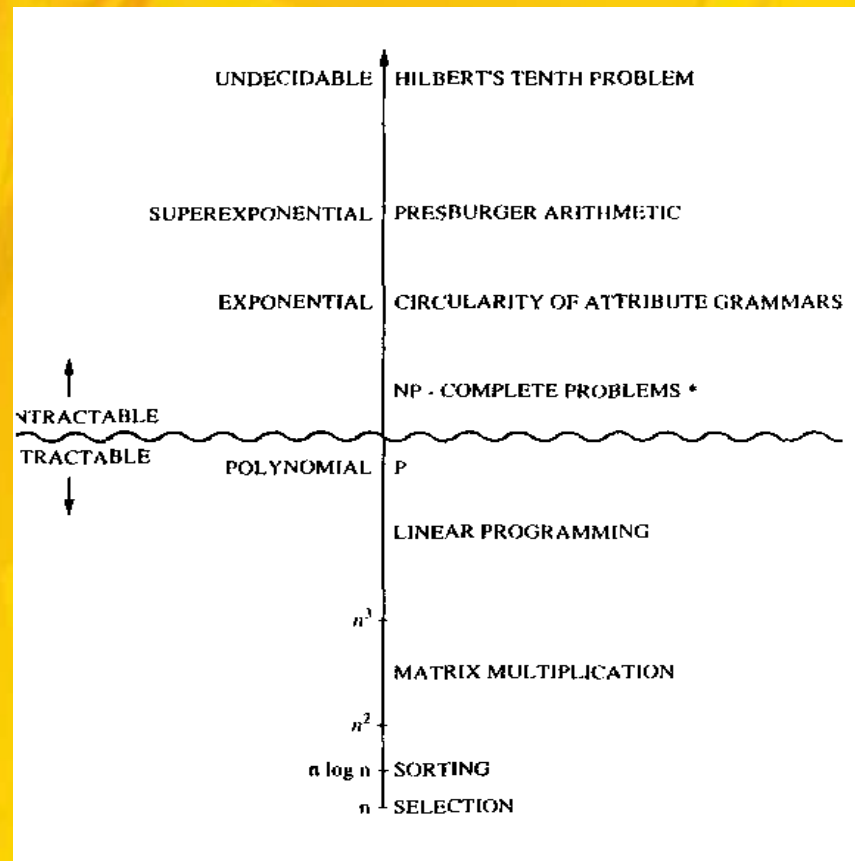


Figure 3: The Spectrum of Computational Complexity

Parallel Performance Limitations

- Intuitively, one might hope that an n processor implementation of an algorithm or parallelized program will achieve the result in $1/n$ (or less) of the time required by a single processor working on the same problem.
- Unfortunately, this is at best an upper bound on the achievable performance.

Metrics for Speedup

For a parallel implementation, *speedup* is the ratio:

$$speedup = \frac{\text{processing time with single processor}}{\text{processing time with } n \text{ processors}} = \frac{t_s}{t_p} \quad (1)$$

On this basis, the theoretically achievable maximum speedup for an n -processor implementation of an algorithm is n .

Amdahl's Law

Amdahl's law explicitly predicts speedup as a function of the fraction of required serial computations in a given algorithm.

Define the following quantities:

s : the fraction of computations which are necessarily serial;

p : the fraction of computations which may be done in parallel
(note $s + p = 1$); and

n : number of processors over which p is distributed.

Amdahl's law predicts speedup as:

$$speedup = \frac{(s + p)}{(s + p/n)} = \frac{1}{(s + p/n)} \quad (2)$$

In general, Amdahl's law is a pessimistic bound on achievable speedup, even for small values of s ($< 5\%$).